

Deploying NGINX and NGINX Plus on Docker

Deploy NGINX and NGINX Plus as the Docker container.

NGINX Plus, the high-performance application delivery platform, load balancer, and web server, is available as the Docker container.

Prerequisites

- Docker installation
- Docker Hub account (NGINX Open Source)
- *nginx-repo.crt* and *nginx-repo.key* files, Dockerfile for Docker image creation (NGINX Plus)

Running NGINX Open Source in a Docker Container

You can create an NGINX instance in a Docker container using the NGINX Open Source image from the Docker Hub.

1. Launch an instance of NGINX running in a container and using the default NGINX configuration with the following command:

```
$ docker run --name mynginx1 -p 80:80 -d nginx
```

where:

- *mynginx1* is the name of the created container based on the NGINX image
- the *-d* option specifies that the container runs in detached mode: the container continues to run until stopped but does not respond to commands run on the command line.
- the *-p* option tells Docker to map the ports exposed in the container by the NGINX image (port 80) to the specified port on the Docker host. The first parameter specifies the port in the Docker host, the second parameter is mapped to the port exposed in the container

The command returns the long form of the container ID:

fcd1fb01b14557c7c9d991238f2558ae2704d129cf9fb97bb4fadf673a58580d. This form of ID is used in the name of log files.

2. Verify that the container was created and is running with the `docker ps` command:

```
$ docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        ...
fcd1fb01b145   nginx:latest   "nginx -g 'daemon of   16 seconds ago Up 15 seconds ...

... PORTS          NAMES
... 0.0.0.0:80->80/tcp mynginx1
```

This command also allows viewing the port mappings set in the previous step: the `PORTS` field in the output reports that port 80 on the Docker host is mapped to port 80 in the container.

Running NGINX Plus in a Docker Container

Docker can also be used with NGINX Plus. The difference between using Docker with NGINX Open Source is that you first need to create an NGINX Plus image, because as a commercial offering NGINX Plus is not available at Docker Hub.

Note: Never upload your NGINX Plus images to a public repository such as Docker Hub. Doing so violates your license agreement.

Creating NGINX Plus Docker Image

To generate an NGINX Plus image:

1. Create the Docker build context, or a Dockerfile:

```
1 FROM debian:bullseye-slim
2
3 LABEL maintainer="NGINX Docker Maintainers <docker-maint@nginx.com>"
4
5 # Define NGINX versions for NGINX Plus and NGINX Plus modules
6 # Uncomment this block and the versioned nginxPackages block in the main RUN
7 # instruction to install a specific release
8 # ENV NGINX_VERSION 29
9 # ENV NJS_VERSION 0.7.12
10 # ENV PKG_RELEASE 1~bullseye
```

```

11
12 # Download certificate and key from the customer portal (https://account.f5.com)
13 # and copy to the build context
14 RUN --mount=type=secret,id=nginx-crt,dst=nginx-repo.crt \
15     --mount=type=secret,id=nginx-key,dst=nginx-repo.key \
16     set -x \
17 # Create nginx user/group first, to be consistent throughout Docker variants
18     && addgroup --system --gid 101 nginx \
19     && adduser --system --disabled-login --ingroup nginx --no-create-home --home /nonexistent --gecos "nginx user" --shell /bin/false --u
20     && apt-get update \
21     && apt-get install --no-install-recommends --no-install-suggests -y \
22         ca-certificates \
23         gnupg1 \
24         lsb-release \
25     && \
26     NGINX_GPGKEY=573BFD6B3D8FBC641079A6ABABF5BD827BD9BF62; \
27     NGINX_GPGKEY_PATH=/usr/share/keyrings/nginx-archive-keyring.gpg; \
28     export GNUPGHOME="$(mktemp -d)"; \
29     found=''; \
30     for server in \
31         hkp://keyserver.ubuntu.com:80 \
32         pgp.mit.edu \
33     ; do \
34         echo "Fetching GPG key $NGINX_GPGKEY from $server"; \
35         gpg1 --keyserver "$server" --keyserver-options timeout=10 --recv-keys "$NGINX_GPGKEY" && found=yes && break; \
36     done; \
37     test -z "$found" && echo >&2 "error: failed to fetch GPG key $NGINX_GPGKEY" && exit 1; \
38     gpg1 --export "$NGINX_GPGKEY" > "$NGINX_GPGKEY_PATH" ; \
39     rm -rf "$GNUPGHOME"; \
40     apt-get remove --purge --auto-remove -y gnupg1 && rm -rf /var/lib/apt/lists/* \
41 # Install the latest release of NGINX Plus and/or NGINX Plus modules
42 # Uncomment individual modules if necessary
43 # Use versioned packages over defaults to specify a release
44     && nginxPackages=" \
45         nginx-plus \
46         # nginx-plus=${NGINX_VERSION}-${PKG_RELEASE} \
47         # nginx-plus-module-xslt \
48         # nginx-plus-module-xslt=${NGINX_VERSION}-${PKG_RELEASE} \
49         # nginx-plus-module-geoip \
50         # nginx-plus-module-geoip=${NGINX_VERSION}-${PKG_RELEASE} \
51         # nginx-plus-module-image-filter \
52         # nginx-plus-module-image-filter=${NGINX_VERSION}-${PKG_RELEASE} \
53         # nginx-plus-module-perl \
54         # nginx-plus-module-perl=${NGINX_VERSION}-${PKG_RELEASE} \
55         # nginx-plus-module-njs \
56         # nginx-plus-module-njs=${NGINX_VERSION}+${NJS_VERSION}-${PKG_RELEASE} \
57     " \
58     && echo "Acquire::https::pkgs.nginx.com::Verify-Peer \"true\";" > /etc/apt/apt.conf.d/90nginx \
59     && echo "Acquire::https::pkgs.nginx.com::Verify-Host \"true\";" >> /etc/apt/apt.conf.d/90nginx \
60     && echo "Acquire::https::pkgs.nginx.com::SslCert \"/etc/ssl/nginx/nginx-repo.crt\";" >> /etc/apt/apt.conf.d/90nginx \
61     && echo "Acquire::https::pkgs.nginx.com::SslKey \"/etc/ssl/nginx/nginx-repo.key\";" >> /etc/apt/apt.conf.d/90nginx \
62     && printf "deb [signed-by=/usr/share/keyrings/nginx-archive-keyring.gpg] https://pkgs.nginx.com/plus/debian `lsb_release -cs` nginx-p
63     && mkdir -p /etc/ssl/nginx \
64     && cat nginx-repo.crt > /etc/ssl/nginx/nginx-repo.crt \
65     && cat nginx-repo.key > /etc/ssl/nginx/nginx-repo.key \
66     && apt-get update \
67     && apt-get install --no-install-recommends --no-install-suggests -y \
68         $nginxPackages \
69         curl \
70         gettext-base \
71     && apt-get remove --purge -y lsb-release \
72     && apt-get remove --purge --auto-remove -y && rm -rf /var/lib/apt/lists/* /etc/apt/sources.list.d/nginx-plus.list \
73     && rm -rf /etc/apt/apt.conf.d/90nginx /etc/ssl/nginx \
74 # Forward request logs to Docker log collector
75     && ln -sf /dev/stdout /var/log/nginx/access.log \
76     && ln -sf /dev/stderr /var/log/nginx/error.log
77
78 EXPOSE 80
79

```

```
80  STOPSIGNAL SIGQUIT
81
82  CMD ["nginx", "-g", "daemon off;"]
```

Dockerfile hosted with ❤️ by GitHub

view raw

```
1  FROM alpine:3.17
2
3  LABEL maintainer="NGINX Docker Maintainers <docker-maint@nginx.com>"
4
5  # Define NGINX versions for NGINX Plus and NGINX Plus modules
6  # Uncomment this block and the versioned nginxPackages in the main RUN
7  # instruction to install a specific release
8  # ENV NGINX_VERSION 29
9  # ENV NJS_VERSION    0.7.12
10 # ENV PKG_RELEASE    1
11
12 # Download certificate and key from the customer portal (https://account.f5.com)
13 # and copy to the build context
14 RUN --mount=type=secret,id=nginx-crt,dst=cert.pem \
15     --mount=type=secret,id=nginx-key,dst=cert.key \
16     set -x \
17 # Create nginx user/group first, to be consistent throughout Docker variants
18     && addgroup -g 101 -S nginx \
19     && adduser -S -D -H -u 101 -h /var/cache/nginx -s /sbin/nologin -G nginx -g nginx nginx \
20 # Install the latest release of NGINX Plus and/or NGINX Plus modules
21 # Uncomment individual modules if necessary
22 # Use versioned packages over defaults to specify a release
23     && nginxPackages=" \
24         nginx-plus \
25         # nginx-plus=${NGINX_VERSION}-r${PKG_RELEASE} \
26         # nginx-plus-module-xslt \
27         # nginx-plus-module-xslt=${NGINX_VERSION}-r${PKG_RELEASE} \
28         # nginx-plus-module-geoip \
29         # nginx-plus-module-geoip=${NGINX_VERSION}-r${PKG_RELEASE} \
30         # nginx-plus-module-image-filter \
31         # nginx-plus-module-image-filter=${NGINX_VERSION}-r${PKG_RELEASE} \
32         # nginx-plus-module-perl \
33         # nginx-plus-module-perl=${NGINX_VERSION}-r${PKG_RELEASE} \
34         # nginx-plus-module-njs \
35         # nginx-plus-module-njs=${NGINX_VERSION}.${NJS_VERSION}-r${PKG_RELEASE} \
36     " \
37     KEY_SHA512="e09fa32f0a0eab2b879ccbbc4d0e4fb9751486eedda75e35fac65802cc9faa266425edf83e261137a2f4d16281ce2c1a5f4502930fe75154723da0142
38     && wget -O /tmp/nginx_signing.rsa.pub https://nginx.org/keys/nginx_signing.rsa.pub \
39     && if echo "$KEY_SHA512 */tmp/nginx_signing.rsa.pub" | sha512sum -c -; then \
40         echo "key verification succeeded!"; \
41         mv /tmp/nginx_signing.rsa.pub /etc/apk/keys/; \
42     else \
43         echo "key verification failed!"; \
44         exit 1; \
45     fi \
46     && cat cert.pem > /etc/apk/cert.pem \
47     && cat cert.key > /etc/apk/cert.key \
48     && apk add -X "https://pkgs.nginx.com/plus/alpine/v$(egrep -o '^[0-9]+\.[0-9]+' /etc/alpine-release)/main" --no-cache $nginxPackages
49     && if [ -f "/etc/apk/keys/nginx_signing.rsa.pub" ]; then rm -f /etc/apk/keys/nginx_signing.rsa.pub; fi \
50     && if [ -f "/etc/apk/cert.key" ] && [ -f "/etc/apk/cert.pem" ]; then rm -f /etc/apk/cert.key /etc/apk/cert.pem; fi \
51 # Bring in tzdata so users could set the timezones through the environment
52 # variables
53     && apk add --no-cache tzdata \
54 # Bring in curl and ca-certificates to make registering on DNS SD easier
55     && apk add --no-cache curl ca-certificates \
56 # Forward request and error logs to Docker log collector
57     && ln -sf /dev/stdout /var/log/nginx/access.log \
58     && ln -sf /dev/stderr /var/log/nginx/error.log
59
60 EXPOSE 80
61
62 STOPSIGNAL SIGQUIT
```

```
63
64  CMD ["nginx", "-g", "daemon off;"]
65
66  # vim:syntax=Dockerfile
```

Dockerfile.alpine hosted with ❤️ by GitHub

view raw

2. As with NGINX Open Source, default NGINX Plus image has the same default settings:
- access and error logs are linked to the Docker log collector
 - no volumes are specified: a Dockerfile can be used to create base images from which you can create new images with volumes specified, or volumes can be specified manually:

```
VOLUME /usr/share/nginx/html
VOLUME /etc/nginx
```

- no files are copied from the Docker host as a container is created: you can add `COPY` definitions to each Dockerfile, or the image you create can be used as the basis for another image
3. Log in to [MyF5 Customer Portal](#) and download your *nginx-repo.crt* and *nginx-repo.key* files. For a trial of NGINX Plus, the files are provided with your trial package.
4. Copy the files to the directory where the Dockerfile is located.
5. Create a Docker image, for example, `nginxplus` (note the final period in the command).

```
$ docker build --no-cache --secret id=nginx-key,src=nginx-repo.key --secret id=nginx-crt,src=nginx-repo.crt -t
```

The `--no-cache` option tells Docker to build the image from scratch and ensures the installation of the latest version of NGINX Plus. If the Dockerfile was previously used to build an image without the `--no-cache` option, the new image uses the version of NGINX Plus from the previously built image from the Docker cache.

6. Verify that the `nginxplus` image was created successfully with the `docker images` command:

```
$ docker images nginxplus
REPOSITORY TAG IMAGE ID CREATED SIZE
nginxplus latest ef2bf65931cf 6 seconds ago 91.2 MB
```

7. Create a container based on this image, for example, `mynginxplus` container:

```
$ docker run --name mynginxplus -p 80:80 -d nginxplus
```

8. Verify that the `mynginxplus` container is up and running with the `docker ps` command:

```
$ docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS
eb7be9f439db nginxplus:latest "nginx -g 'daemon of 1 minute ago Up 15 seconds ...

... PORTS NAMES
... 0.0.0.0:80->80/tcp mynginxplus
```

NGINX Plus containers are controlled and managed in the same way as NGINX Open Source containers.

Managing Content and Configuration Files

Content served by NGINX and NGINX configuration files can be managed in several ways:

- files are maintained on the Docker host
- files are copied from the Docker host to a container
- files are maintained in the container

Maintaining Content and Configuration Files on the Docker Host

When the container is created, you can mount a local directory on the Docker host to a directory in the container. The NGINX image uses the default NGINX configuration, which uses `/usr/share/nginx/html` as the container's root directory and puts configuration files in `/etc/nginx`. For a Docker host with content in the local directory `/var/www` and configuration files in `/var/nginx/conf`, run the command:

```
$ docker run --name mynginx2 \
  --mount type=bind,source=/var/www,target=/usr/share/nginx/html,readonly \
  --mount type=bind,source=/var/nginx/conf,target=/etc/nginx/conf,readonly \
  -p 80:80 \
  -d nginxplus
```

Any change made to the files in the local directories `/var/www` and `/var/nginx/conf` on the Docker host are reflected in the directories `/usr/share/nginx/html` and `/etc/nginx` in the container. The `readonly` option means these directories can be changed only on the Docker host, not from within the container.

Copying Content and Configuration Files from the Docker Host

Docker can copy the content and configuration files from a local directory on the Docker host during container creation. Once a container is created, the files are maintained by creating a new container when files change or by modifying the files in the container.

A simple way to copy the files is to create a Dockerfile with commands that are run during generation of a new Docker image based on the NGINX image. For the file-copy (COPY) commands in the Dockerfile, the local directory path is relative to the build context where the Dockerfile is located.

Let's assume that the content directory is `content` and the directory for configuration files is `conf`, both subdirectories of the directory where the Dockerfile is located. The NGINX image has the default NGINX configuration files, including `default.conf`, in the `/etc/nginx/conf.d` directory. To use the configuration files from the Docker host only, delete the default files with the `RUN` command:

```
FROM nginx
RUN rm /etc/nginx/conf.d/default.conf
COPY content /usr/share/nginx/html
COPY conf /etc/nginx
```

Create NGINX image by running the command from the directory where the Dockerfile is located. The period (".") at the end of the command defines the current directory as the build context, which contains the Dockerfile and the directories to be copied:

```
$ docker build -t mynginx_image1 .
```

Create a container `mynginx3` based on the `mynginx_image1` image:

```
$ docker run --name mynginx3 -p 80:80 -d mynginx_image1
```

To make changes to the files in the container, use a helper container as described in the next section.

Maintaining Content and Configuration Files in the Container

As SSH cannot be used to access the NGINX container, to edit the content or configuration files directly you need to create a helper container that has shell access. For the helper container to have access to the files, create a new image that has the proper Docker data volumes defined for the image:

1. Copy nginx content and configuration files and define the volume for the image with the Dockerfile:

```
FROM nginx
COPY content /usr/share/nginx/html
COPY conf /etc/nginx
VOLUME /usr/share/nginx/html
VOLUME /etc/nginx
```

2. Create the new NGINX image by running the following command:

```
$ docker build -t mynginx_image2 .
```

3. Create an NGINX container `mynginx4` based on the `mynginx_image2` image:

```
$ docker run --name mynginx4 -p 80:80 -d mynginx_image2
```

4. Start a helper container `mynginx4_files` that has a shell, providing access the content and configuration directories of the `mynginx4` container we just created:

```
$ docker run -i -t --volumes-from mynginx4 --name mynginx4_files debian /bin/bash
root@b1cbbad63dd1:/#
```

where:

- the new `mynginx4_files` helper container runs in the foreground with a persistent standard input (the `-i` option) and a tty (the `-t` option). All volumes defined in `mynginx4` are mounted as local directories in the helper container.

- the `debian` argument means that the helper container uses the Debian image from Docker Hub. Because the NGINX image also uses Debian, it is most efficient to use Debian for the helper container, rather than having Docker load another operating system
- the `/bin/bash` argument means that the bash shell runs in the helper container, presenting a shell prompt that you can use to modify files as needed

To start and stop the container, run the commands:

```
$ docker start mynginx4_files
$ docker stop mynginx4_files
```

To exit the shell but leave the container running, press `Ctrl+p` followed by `Ctrl+q`. To regain shell access to a running container, run this command:

```
$ docker attach mynginx4_files
```

To exit the shell and terminate the container, run the `exit` command.

Managing Logging

You can use default logging or customize logging.

Using Default Logging

By default, the NGINX image is configured to send NGINX [access log](#) and [error log](#) to the Docker log collector. This is done by linking them to `stdout` and `stderr`: all messages from both logs are then written to the file

`/var/lib/docker/containers/container-ID/container-ID-json.log` on the Docker host. The container-ID is the long-form ID returned when you [create a container](#). To display the long form ID, run the command:

```
$ docker inspect --format '{{.Id}}' container-name
```

You can use both the Docker command line and the Docker Engine API to extract the log messages.

To extract log messages from the command line, run the command:

```
$ docker logs container-name
```

To extract log messages using the Docker Remote API, send a `GET` request using the Docker Unix sock:

```
$ curl --unix-sock /var/run/docker-sock http://localhost/containers/container-name/logs?stdout=1&stderr=1
```

To include only access log messages in the output, include only `stdout=1`. To limit the output to error log messages, include only `stderr=1`. For other available options, see [Get container logs](#) section of the [Docker Engine API](#) documentation.

Using Customized Logging

If you want to configure logging differently for certain configuration blocks (such as `server {}` and `location {}`), define a Docker volume for the directory in which to store the log files in the container, create a helper container to access the log files, and use any logging tools. To implement this, create a new image that contains the volume or volumes for the logging files.

For example, to configure NGINX to store log files in `/var/log/nginx/log`, add a `VOLUME` definition for this directory to the Dockerfile (provided that content and configuration Files are [managed in the container](#)):

```
FROM nginx
COPY content /usr/share/nginx/html
COPY conf /etc/nginx
VOLUME /var/log/nginx/log
```

Then you can [create an image](#) and use it to create an NGINX container and a helper container that have access to the logging directory. The helper container can have any desired logging tools installed.

Controlling NGINX

Since there is no direct access to the command line of the NGINX container, NGINX commands cannot be sent to a container directly. Instead, [signals](#) can be sent to a container via Docker `kill` command.

To reload the NGINX configuration, send the `HUP` signal to Docker:

```
$ docker kill -s HUP container-name
```

To restart NGINX, run this command to restart the container:

```
$ docker restart container-name
```

